

Intro: Coin Changing

Coin Changing

Goal. Given currency denominations: 1, 5, 10, 25, 100, give change to customer using *fewest* number of coins.

Ex: 34¢.



Algorithm is
“Greedy”: One
large coin
better than two
or more
smaller ones

Cashier's algorithm. At each iteration, give the *largest* coin valued \leq the amount to be paid.

Ex: \$2.89.



Coin-Changing: Does Greedy Always Work?


Observation. Greedy algorithm is sub-optimal for US postal denominations: 1, 10, 21, 34, 70, 100, 350, 1225, 1500.

Counterexample. 140¢.

- Greedy: 100, 34, 1, 1, 1, 1, 1, 1.
- Optimal: 70, 70.

Algorithm is “Greedy”,
but also short-sighted
– attractive choice
now may lead to dead
ends later.

Correctness is key!



Outline & Goals

“Greedy Algorithms”

what they are

Pros

intuitive

often simple

often fast

Cons

often incorrect!

Proof techniques

stay ahead

structural

exchange arguments


4.1 Interval Scheduling

Proof Technique 1: “greedy stays ahead”

Interval Scheduling

Interval scheduling.

- Job j starts at s_j and finishes at f_j .
- Two jobs **compatible** if they don't overlap.
- Goal: find maximum subset of mutually compatible jobs.



Interval Scheduling: Greedy Algorithms

Greedy template. Consider jobs in some order. Take each job provided it's compatible with the ones already taken.

- What order?
- Does that give best answer?
- Why or why not?
- Does it help to be greedy about order?

Interval Scheduling: Greedy Algorithms

Greedy template. Consider jobs in some order. Take each job provided it's compatible with the ones already taken.

[Earliest start time] Consider jobs in ascending order of start time s_j .

[Earliest finish time] Consider jobs in ascending order of finish time f_j .

[Shortest interval] Consider jobs in ascending order of interval length $f_j - s_j$.

[Fewest conflicts] For each job, count the number of conflicting jobs c_j .
Schedule in ascending order of conflicts c_j .

Interval Scheduling: Greedy Algorithm



Greedy algorithm. Consider jobs in increasing order of finish time. Take each job provided it's compatible with the ones already taken.

```
Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ .  
  ↙ jobs selected  
A = {}  
for j = 1 to n {  
    if (job j compatible with A)  
        A = A ∪ {j}  
}  
return A
```




Implementation. $O(n \log n)$.

- Remember job j^* that was added last to A.
- Job j is compatible with A if $s_j \geq f_{j^*}$.



Interval Scheduling





Interval Scheduling





Interval Scheduling





Interval Scheduling





Interval Scheduling





Interval Scheduling





Interval Scheduling



Interval Scheduling



Interval Scheduling



Interval Scheduling: Correctness

Theorem. Greedy algorithm is optimal.

Pf. (“greedy stays ahead”)


Let i_1, i_2, \dots, i_k be jobs picked by greedy, j_1, j_2, \dots, j_m those in some optimal solution

Show $f(i_r) \leq f(j_r)$ by induction on r .

Basis: i_1 chosen to have min finish time, so $f(i_1) \leq f(j_1)$

Ind: $f(i_r) \leq f(j_r) \leq s(j_{r+1})$, so j_{r+1} is among the candidates considered by greedy when it picked i_{r+1} , & it picks min finish, so $f(i_{r+1}) \leq f(j_{r+1})$

Similarly, $k \geq m$, else j_{k+1} is among (nonempty) set of candidates for i_{k+1}



4.1 Interval Partitioning


Proof Technique 2: “Structural”

Interval Partitioning

Interval partitioning.

- Lecture j starts at s_j and finishes at f_j .
- Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

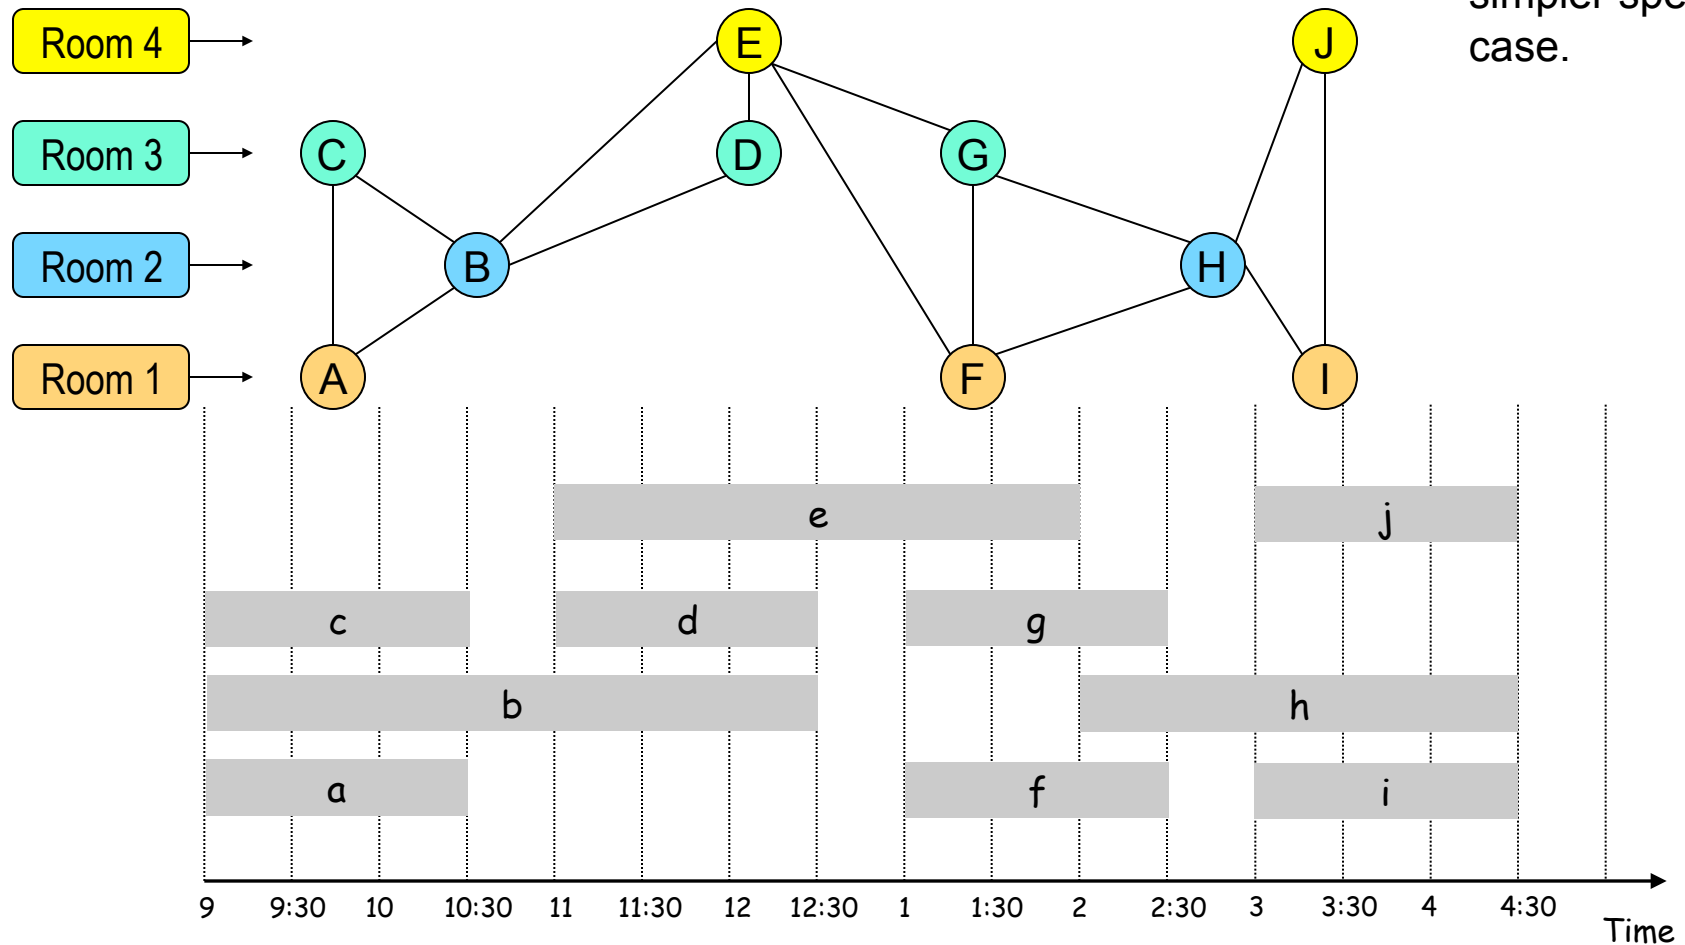
Ex: This schedule uses 4 classrooms to schedule 10 lectures.



Interval Partitioning as Interval Graph Coloring

Vertices = classes;
edges = conflicting class pairs;
different colors = different assigned rooms

Note: graph coloring is very hard in general, but graphs corresponding to interval intersections are a much simpler special case.




Interval Partitioning

Interval partitioning.

- Lecture j starts at s_j and finishes at f_j .
- Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Ex: This schedule uses only 3.



Interval Partitioning: A “Structural” Lower Bound on Optimal Solution

Def. The depth of a set of open intervals is the maximum number that contain any given time.


↑
no collisions at ends

Key observation. Number of classrooms needed \geq depth.

Ex: Depth of schedule below = 3 \Rightarrow schedule below is optimal.

↑
a, b, c all contain 9:30

Q. Does there always exist a schedule equal to depth of intervals?




Interval Partitioning: Greedy Algorithm

Greedy algorithm. Consider lectures in increasing order of start time: assign lecture to any compatible classroom.


```
Sort intervals by starting time so that  $s_1 \leq s_2 \leq \dots \leq s_n$ .  
d = 0      ← number of allocated classrooms  
  
for j = 1 to n {  
    if (lect j is compatible with some classroom k,  $1 \leq k \leq d$ )  
        schedule lecture j in classroom k  
    else  
        allocate a new classroom d + 1  
        schedule lecture j in classroom d + 1  
        d = d + 1  
}
```

Implementation? Run-time?
Exercises


Interval Partitioning




Interval Partitioning




Interval Partitioning




Interval Partitioning




Interval Partitioning




Interval Partitioning




Interval Partitioning




Interval Partitioning




Interval Partitioning



Interval Partitioning



Interval Partitioning



Interval Partitioning: Greedy Analysis

Theorem. Greedy algorithm is optimal.

Pf (exploit structural property).

- Let d = number of classrooms that the greedy algorithm allocates.
- Classroom d is opened because we needed to schedule a job, say j , that is incompatible with all $d-1$ previously used classrooms.
- Since we sorted by start time, all these incompatibilities are caused by lectures that start no later than s_j .
- Thus, we have d lectures overlapping at time s_j , i.e. $\text{depth} \geq d$
- “Key observation” all schedules use $\geq \text{depth}$ classrooms, so $d = \text{depth}$ and greedy is optimal ▪